Technische Hochschule Deggendorf

MDM-M-2 Case Study Intelligent Systems in Production

Binary Classification of AOI Alarms into True Defects and False Calls

Second check of AOI results with supervised machine learning
(Siemens production data)

Bibliography Report

Team 1

Yu Han
Melisa Subasi
Meron Hagos
Alejandro Rivas

Under the supervision of:
Prof. Dr.-Ing. Hamidreza Heidari

Date of submission: 03.11.2026

Abstract

In electronics assembly, automated optical inspection (AOI) checks every solder joint after soldering and sends every flagged board to a manual inspection station. Most of these alarms are false calls: in public Siemens production data, about 99 % of the flagged solder joints are false calls. A second check with supervised machine learning after the AOI could release many false calls without manual inspection, and Siemens is working on this problem in its Erlangen electronics factory today.

This report reviews 18 sources, from the general to the specific. Earlier studies show that classifiers such as random forests can recognise false calls from AOI measurements. On the same Siemens data, however, models that meet the targets on a test set from the same period (at most 1 % missed defects, at least 40 % fewer manual checks) miss 5–15 % of the true defects on every later production period. The literature on dataset shift and concept drift, cases from manufacturing, and the authors of the data and of the study themselves point to drift, a change of the data over time, as the cause; an unstable method for setting the decision threshold adds to it. The standard remedy in the drift literature is to retrain the model regularly on recent data. In related applications, a binary fraud classifier with less than 1 % positives and a quality model on a semiconductor materials line, regular retraining reduced the loss in performance. For the second check of AOI results, this remedy has not been tested in the literature we reviewed.

Our project therefore reimplements the method of Pfab and Rothering with its own code, on the same data with the same models, and addresses the two causes they name: the threshold method and retraining. A design of experiments varies the threshold method (3 levels), the retraining schedule (6 levels), the retraining data (2 levels) and the model (2 levels). The deliverable is a simulation of the second check in production in which threshold and retraining are settings; its results are compared with the published result, and the targets of the original study (at most 1 % missed true defects, at least 40 % fewer manual checks) serve as reference lines.

Contents

Abstract i

1 Introduction 1

1.1 Motivation . . . 1

1.2 Key terms . . . 3

2 Materials and Methods 3

2.1 Reviewed literature . . . 3

2.2 Summary of the reviewed works . . . 8

3 Discussion 8

3.1 Why the earlier results look better than they are . . . 8

3.2 Limitations of the reviewed sources . . . 9

4 Conclusion 9

References 10

List of Figures

1 The reviewed literature from general to specific. Each box is one level of the review, with its sources on the right; Section 2 follows them from top to bottom. The line below the funnel is our project. . . . 1

2 Missed true defects (horizontal) and false calls released (vertical) of the two models of [7], on the test set (stars, [7, p. 11, Table V]) and on the five later periods (dots, [7, p. 7, Table IV]). Further right is worse, higher is better. Dashed lines: the two limits of [7], at most 1 % missed (vertical) and at least 40 % released (horizontal); the green area between them is the target zone. . . . 2

3 Share of flagged boards with a true defect per production week (bars, left axis; red: weeks above 3 %) and number of flagged boards per week (line, right axis). Data of [4], our analysis. . . . 5

4 Missed true defects of the two models of [7] on the test set and on the five later periods (slices), mean and standard deviation of their runs [7, p. 7, Table IV; p. 11, Table V]. Dashed line: the target of at most 1 %. . . . 6

5 Minimum number of true defects in the left-out data so that the threshold keeps the share of missed true defects below α with the stated confidence, $n \geq \log \delta / \log(1-\alpha)$ [14]; dark bars 95 %, light bars 90 % confidence; computed by us. . . . 7

List of Tables

1 What a second check after AOI has reached so far. About 99 % of the flagged joints are false calls, so the share of false calls released is roughly the share of manual checks saved. 2

2 Key terms used in this report. 3

3 The 18 reviewed sources in the order of Figure 1: the finding we use and where it is in the source. 4

4 Studies on AOI false calls compared with our project. 8

5 Studies on drift and retraining compared with our project. 8

6 From the literature to the variables of our simulation (the two next steps named in [7]). . 9

1 Introduction

Zero-defect manufacturing relies on detecting defects before a product leaves the factory [1]. In electronics assembly, components are soldered onto printed circuit boards by surface-mount technology (SMT), and AOI is placed after soldering [2]. Every board that the AOI flags goes to a manual inspection station, where an operator decides whether it is a true defect or a false call [4]. A classifier after the AOI, a second check, could release the false calls without a manual check. This review covers 18 sources from 2001 to 2026 on this second check; image-based defect detection is outside its scope, because the second check starts after the AOI has made its measurements. Figure 1 shows the 18 sources from the general goal to our project; Section 2 follows it level by level.

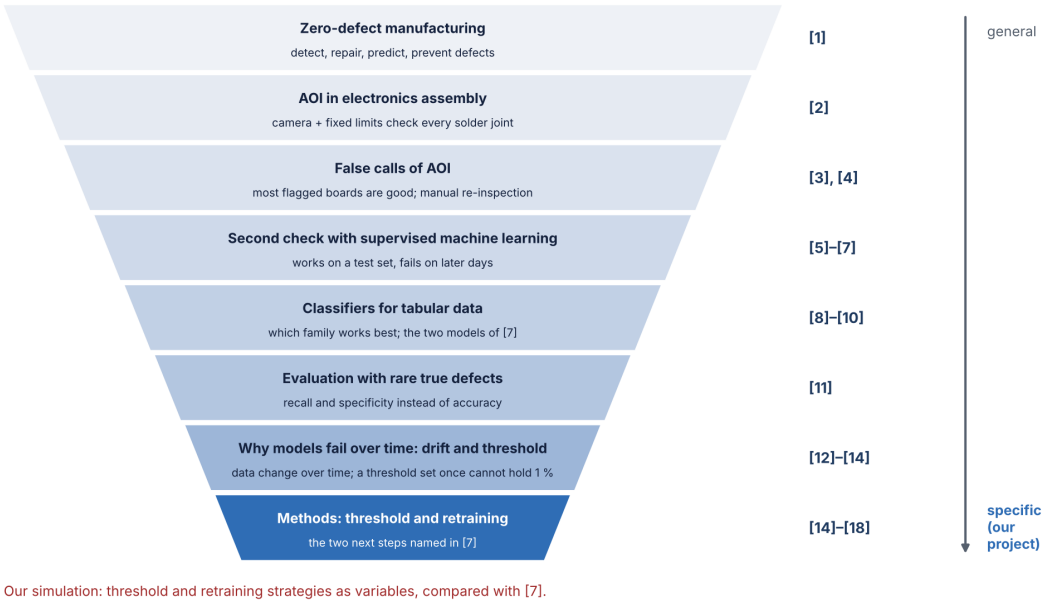


Figure 1: The reviewed literature from general to specific. Each box is one level of the review, with its sources on the right; Section 2 follows them from top to bottom. The line below the funnel is our project.

1.1 Motivation

Every board that the AOI flags is checked by a person, and about 99 % of the flagged solder joints are false calls [4, 7]. This costs operator time on every line, and it is a current problem: Siemens is working on an AI second check in its Erlangen factory today, with the target of 40 % fewer false calls [3]. Earlier studies show that a classifier on the AOI measurements can recognise many false calls [5, 6]. The only study that tests such a classifier on later production periods, however, finds that it fails there: it misses 5–15 % of the true defects instead of at most 1 % [7]. A factory cannot use a second check that lets defects through. Table 1 shows what has been reached so far and what is at stake.

Table 1: What a second check after AOI has reached so far. About 99 % of the flagged joints are false calls, so the share of false calls released is roughly the share of manual checks saved.

Source	Setting	False calls released (manual checks saved)	Missed true defects
Siemens [3]	industrial target, Erlangen	40 % (target)	–
Pfab & Rothering [7, p. 6, Table III]	test set, ideal threshold (best possible)	92.5 % (XGBoost), 90.1 % (RF), 79.6 % (BRF)	1 % (by construction)
Pfab & Rothering [7, p. 11, Table V]	test set, threshold as set by their method	79.7 % (BRF), 89.7 % (RF)	0.8 % (BRF), 1.0 % (RF)
Pfab & Rothering [7, p. 7, Table IV]	five later production periods	41.5–78.0 % (BRF), 39.7–76.0 % (RF)	5.3–9.4 % (BRF), 10.2–15.0 % (RF)
Our baseline	five later periods, 10 seeds, mean	57.9 % (BRF), 55.8 % (RF)	7.5 % (BRF), 9.1 % (RF)
Target of [7]	any period	≥ 40 %	≤ 1 %

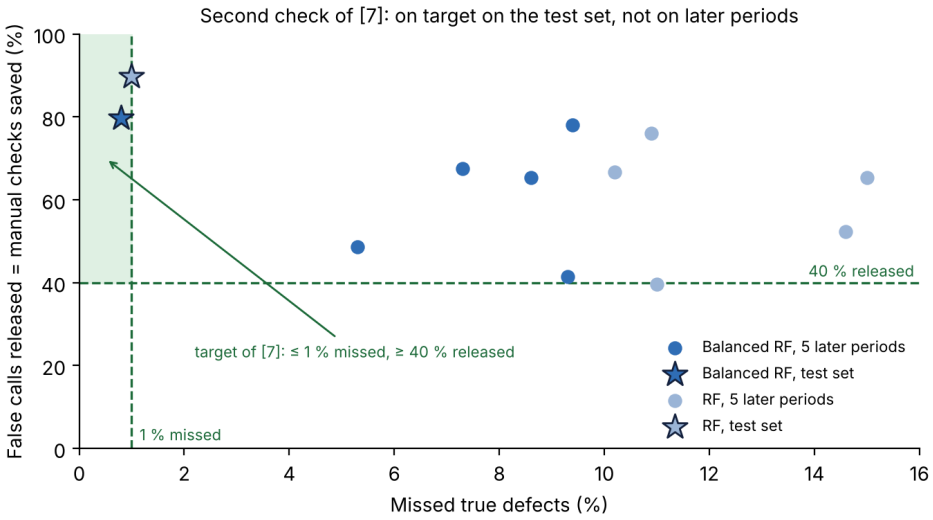


Figure 2: Missed true defects (horizontal) and false calls released (vertical) of the two models of [7], on the test set (stars, [7, p. 11, Table V]) and on the five later periods (dots, [7, p. 7, Table IV]). Further right is worse, higher is better. Dashed lines: the two limits of [7], at most 1 % missed (vertical) and at least 40 % released (horizontal); the green area between them is the target zone.

Table 1 compares the results in the two numbers that matter on the line. Each released false call is one manual check saved; each missed true defect is a defective joint that leaves the station unchecked. The two pull against each other: a lower threshold releases more false calls but also misses more true defects. A result is therefore only meaningful as a pair, and the target of [7] defines the acceptable pairs: at most 1 % missed and at least 40 % released. The 40 % is also the target that Siemens names for its own second check [3].

The second and third rows show what is possible on data from the same period as the training data. With an ideal threshold, the best model, XGBoost, would release 92.5 % of the false calls at the 1 % limit. In the Siemens data, with 440,274 flagged joints of which 4,622 are true defects (1.05 %), this would mean about 400,000 fewer manual joint checks over the 133 days, at the cost of about 46 missed true defects. With the threshold actually set by the method of [7], the random forest comes close to this, 89.7 % released and 1.0 % missed, while XGBoost no longer meets the targets. The gap between the two rows is caused by the threshold alone, since the models are the same.

The fourth row shows what happens on the five later periods. The released share stays high, 40–78 %, so the saving in manual work would still be large. The missed share, however, rises to 5–15 %. Applied to the 4,622 true defects of the data set, this would mean about 250 to 690 missed defective joints instead of about 46. In Figure 2, all ten later results therefore lie to the right of the target area, while the two test-set results lie inside it. Our baseline, the reimplementation of [7], shows the same pattern (fifth row): 55.8–57.9 % released and 7.5–9.1 % missed on average over the later periods.

The table thus locates the problem precisely. The released share already meets the target on every period; what fails is the missed share. A useful fix must bring the missed true defects on later periods back to 1 % without losing the released false calls. If it succeeds, the second check would save at least 40 % of the manual checks, the Siemens target, and up to about 90 % if the test-set result could be kept. The authors of [7] name two such fixes, a better method for setting the threshold and handling drift, but test neither. Our motivation is to find out whether these two fixes make the second check usable in real production.

1.2 Key terms

Table 2 explains the terms used in this report. They follow the core study [7], so that our results can be compared with it.

Table 2: Key terms used in this report.

Term	Meaning
Automated optical inspection (AOI)	A camera system after soldering that measures every solder joint and compares each value with fixed limits [2].
Flagged joint, flagged board	A joint with at least one value outside its limits; a board with at least one flagged joint. Every flagged board goes to a manual inspection station [4].
True defect, false call	The label set at the manual station: the joint is really defective, or the alarm was false. True defect is the positive class.
Second check	A classifier between the AOI and the manual station that decides for each flagged joint whether it can skip the manual check.
Missed true defects (slip)	Share of true defects that the second check releases as false calls, $1 -$ recall. Target in [7]: at most 1 %.
False calls released (volume reduction)	Share of false calls that skip the manual station, the specificity. Target in [7]: at least 40 %.
Threshold	The score above which a joint is treated as a true defect; it is set so that recall reaches 99 % on known data.
Drift	A change of the data over time, so that a model trained on old data no longer fits new data [12].
Retraining	Training the model again on newer labelled data during production.

2 Materials and Methods

2.1 Reviewed literature

The review below follows Figure 1 from top to bottom as one argument. Table 3 lists the sources in the same order, with the finding we use and where it is in the source; page numbers refer to the version linked in the reference list.

Table 3: The 18 reviewed sources in the order of Figure 1: the finding we use and where it is in the source.

Level	Source	Type, year	Finding we use	Where in the source
1	[1]	journal review, 2020	280 articles (1987–2018); four strategies: detection, repair, prediction, prevention	abstract
2	[2]	journal review, 2020	how AOI works: cameras, lighting, feature extraction, fixed limits	abstract
3	[3]	press release, 2026	second check at Erlangen; targets 40 % fewer false calls, 85 % less engineering effort	web page
3	[4]	data article, 2024	132 days, one SMT line; labels set by people; "may contain distribution drifts of unknown type"	abstract; Limitations
4	[5]	conference paper, 2020	kNN, neural network, random forest after AOI; threshold for no extra slip	abstract; summarised in [7, p. 2]
4	[6]	conference paper, 2026	XGBoost best balance of recall and precision on Siemens data; LIME, SHAP	abstract
4	[7]	preprint, 2025	split in time order	p. 3, Sec. IV
			best possible at 1 % missed: 92.5 % (XGBoost), 90.1 % (RF) released	p. 6, Table III
			test set: 0.8 % (BRF) and 1.0 % (RF) missed; 79.7 % and 89.7 % false calls released	p. 11, Table V
			later periods: 5.3–9.4 % (BRF) and 10.2–15.0 % (RF) missed	p. 7, Table IV
			data clusters that "did just appear in a certain period"	p. 3, Fig. 2
			next steps: threshold and distribution drifts; "would fail from the first moment"	pp. 7–8, Sec. VII
5	[8]	conference paper, 2022	45 data sets: tree models remain state of the art on tabular data	p. 1, abstract
5	[9]	journal article, 2001	random forest: trees on random samples and features vote	abstract
5	[10]	technical report, 2004	balanced and weighted random forest for rare classes	pp. 2–3
6	[11]	journal review, 2009	accuracy misleading for rare classes; use recall, specificity	abstract
7	[12]	journal survey, 2014	drift = change of $p(X, y)$; retraining on a sliding window as baseline	definition and adaptation sections
7	[13]	journal article, 2025	controlled drift lowers all metrics of a weld-quality classifier	abstract
7, 8	[14]	journal article, 2018	empirical threshold does not hold an error bound; order statistic, $n \geq \log \delta / \log(1-\alpha)$	p. 2; p. 7 (arXiv copy)
8	[15]	journal article, 2002	re-estimate class shares from unlabelled data and adjust outputs	abstract
8	[16]	conference paper, 2004	drift detection: warning at 2, drift at 3 standard deviations	p. 4, Sec. 3
8	[17]	journal article, 2018	fraud, <1 % positives, 75 million transactions; daily update; recent data better	pp. 1–2; pp. 13–14
8	[18]	preprint, 2026	retraining every 5 vs. 100 batches: about 80 % fewer out-of-range predictions; frequency matters most	p. 9; p. 13

The aim of manufacturing quality is that no defective product leaves the factory. Of the four strategies that research on zero-defect manufacturing distinguishes, detection, repair, prediction and prevention, detection is the one that every line uses today [1]. Detection, however, has a cost: every inspection step adds time and work. A second check after an inspection is one way to lower this cost without letting more defects pass, and this is the idea behind our project.

In electronics assembly, defects are detected after soldering by automated optical inspection. Cameras and lighting record each board, and the AOI software measures every solder joint and compares each value with fixed limits [2]. AOI replaced human visual inspection, which is slow and tires, but it moved the problem rather than solving it. Fixed limits are set conservatively so that no defect slips through, and as a result the machine raises many alarms on joints that are in fact good. Every flagged board still goes to a person.

How large this problem is can be seen in data that Siemens published from one of its own lines. Pfab et al. released 132 days of AOI data from an SMT line of Siemens AG in Germany, with the AOI measurements of every flagged joint and the label from the manual inspection station [4]. About 99 % of the flagged joints are false calls. The authors also warn that the labels are set by people and may be wrong, and that "the data may contain distribution drifts of unknown type". Our own first look at the data confirms the drift (Figure 3): the share of flagged boards with a true defect stays between 0.8 and 2.3 % for 15 weeks and then rises to 3.2–6.3 % in the last four weeks, while the number of flagged boards in these weeks is not unusual. What changes is the mix of true defects and false calls, not the volume, and this will matter later. The problem is also current: in October 2026, Siemens and NVIDIA announced a second inspection step at the Erlangen electronics factory that decides for each AOI alarm whether it is a real defect, with the target of 40 % fewer false calls [3]. This is a target, not a result, but it shows that industry is working on exactly this step.

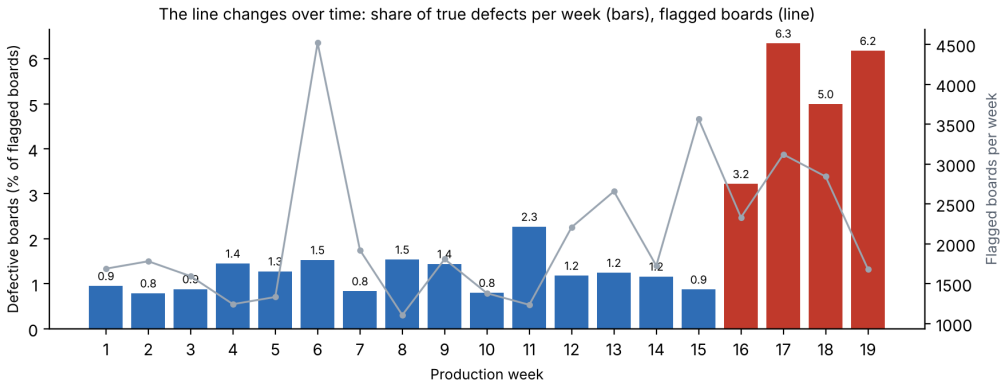


Figure 3: Share of flagged boards with a true defect per production week (bars, left axis; red: weeks above 3 %) and number of flagged boards per week (line, right axis). Data of [4], our analysis.

The natural answer is a second check with supervised machine learning: a classifier that learns from the labels of the manual station and decides for each flagged joint whether it can skip the manual check. Thielen et al. showed that this works in principle on AOI measurements: random forests and neural networks recognised many false calls on large data sets, and the decision threshold could be set so that no additional defects were missed [5]. Albayari et al. reached a similar result on AOI data from Siemens AG, with XGBoost giving the best balance of recall and precision [6]. Neither study reports how many false calls its second check would release at a fixed rate of missed defects, so their results cannot be compared in the terms a factory uses. Both also test the classifier on data from the same period as the training data, by a random split. A factory uses the classifier on the days after training, and neither study tells us what happens then.

Pfab and Rothering answered this question on the public Siemens data, and their study is the core paper of our project [7]. They kept the time order: the first half of the data is used for training and for a test set from the same period, and the second half is cut into five later production periods. Their targets are at most 1 % missed true defects and at least 40 % of the false calls released. On the test set, the best results reported so far for a second check after AOI are theirs. If the threshold is chosen ideally, XGBoost would release 92.5 % of the false calls at the 1 % limit, the random forest 90.1 % and the balanced random forest 79.6 % [7, p. 6, Table III]. With the threshold actually set by their method, XGBoost loses most of this

advantage, and only the random forest and the balanced random forest still meet both targets: they release 89.7 and 79.7 % of the false calls and miss 1.0 and 0.8 % of the true defects [7, p. 11, Table V]. The best model on paper is therefore not the best in use, and the difference is the threshold, a point we return to below. On all five later periods, however, the two models miss 5–15 % of the true defects, already in the first one (Figure 4). The best later result still releases 78 % of the false calls, but misses 9.4 % of the true defects [7, p. 7, Table IV]. The authors conclude that such a model "would fail from the first moment in production". This result changes the question. Whether a second check can recognise false calls is answered; the open question is why it stops working over time, and what can be done about it. Because the authors publish their code and run logs, their result can also be checked and built on directly.

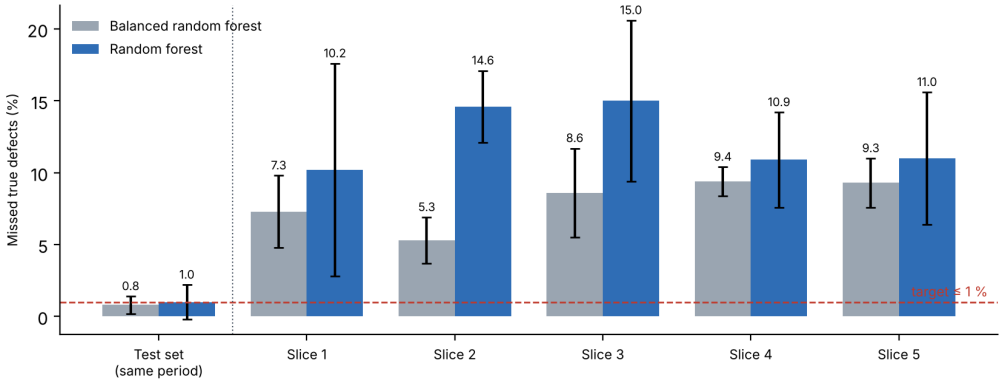


Figure 4: Missed true defects of the two models of [7] on the test set and on the five later periods (slices), mean and standard deviation of their runs [7, p. 7, Table IV; p. 11, Table V]. Dashed line: the target of at most 1 %.

Before turning to the causes, two choices in this study need to be justified, because our project keeps them: the type of model and the way it is measured. AOI measurements are tabular data, one row per joint with numbers and a few categories, and on such data tree ensembles remain stronger than deep learning: on 45 tabular benchmarks, random forests and boosted trees stayed state of the art for medium-sized data [8]. The two models of [7] are of this family. The random forest lets many decision trees, each trained on a random sample of rows and features, vote [9]. The balanced random forest draws the same number of cases from each class for every tree [10], which suits data in which true defects are only about 1 % of the flagged joints. With so few true defects, accuracy is also the wrong measure: a model that calls every joint a false call would be right 99 % of the time and still miss every defect [11]. The study therefore measures, for the two classes separately, the missed true defects and the false calls released, and so do we. The model and the measure are therefore not the reason for the failure; the reason must lie in the data over time.

Pfab and Rothering name two causes themselves. The first is drift. They show data clusters that "did just appear in a certain period of time" and state that "various distribution drifts" occur in the data [7]. In the literature on concept drift, drift is a change of the joint distribution of inputs and labels over time, and a model trained on old data is expected to fail when it occurs [12]. That drift alone is enough to lower a quality classifier in manufacturing was shown in a controlled experiment on friction stir welds, where all metrics dropped once tool geometry, material or lighting changed [13]. Together with the warning of the data authors and the weekly shift in Figure 3, drift is the most likely main cause.

The second cause is the threshold. The threshold of [7] is set once, on the training data, so that 99 % of the known true defects are caught. Tong et al. show why this is not enough [14]: when one type of error must stay below a bound, limiting the error on the data at hand "does not satisfy the [...] error control objective", because the resulting classifiers "are still likely to have [...] errors much larger than" the bound. A threshold set once therefore cannot be trusted to hold 1 % even before the data change, and after a drift it has even less reason to.

Both causes have known remedies, and they are exactly the two next steps that Pfab and Rothering name but do not test: “improving the method for setting the threshold and handling the distribution drifts” [7]. For the threshold, Tong et al. also give a fix. Their method takes the threshold from the scores of a left-out sample in such a way that the error stays below the bound α with a chosen confidence $1-\delta$, a threshold with a safety margin. The price is data: the left-out sample must contain at least $n \geq \log \delta / \log(1-\alpha)$ true defects. For our bound of 1 % with 95 % confidence these are 299 true defects (Figure 5), or roughly 30,000 flagged joints at a defect share of 1 %. The cross-validation folds of [7] are large enough for this, but a short recent period may not be. A second fix addresses the changing share of true defects seen in Figure 3: when the share of a class differs from its share in training, the outputs of a classifier no longer fit, and they can be adjusted by re-estimating the new share from new, even unlabelled data [15]. For a second check this means that the threshold should be set again for each period instead of once.

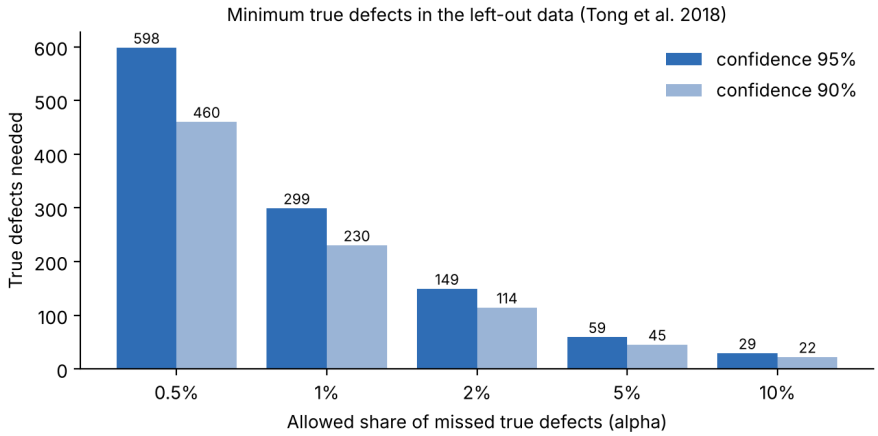


Figure 5: Minimum number of true defects in the left-out data so that the threshold keeps the share of missed true defects below α with the stated confidence, $n \geq \log \delta / \log(1-\alpha)$ [14]; dark bars 95 %, light bars 90 % confidence; computed by us.

For drift, the basic remedy is retraining. In the concept-drift literature, a model that is periodically retrained on a sliding window of the latest data is “often used as a baseline” [12]. The open questions are when to retrain and whether it pays off in practice. One answer to the first question is to retrain only when needed: the drift detection method watches the error rate of the model and retrains when it rises by more than two to three standard deviations above its lowest value [16]. Evidence for the second question comes from two fields close to ours. In credit card fraud detection, a binary task with fewer than 1 % positives on more than 75 million transactions, a classifier updated every day on the newest labelled data performed best, and a model trained on data only a week older was already worse [17]. On a semiconductor materials line, a random forest that predicts product quality degraded after a supplier change and with equipment wear; retraining every 5 batches instead of every 100 gave about 80 % fewer out-of-range predictions, and the frequency of retraining mattered more than the size of the training window [18]. The fraud case is the closest analogue to our task in its statistics, the semiconductor case in its setting; neither is AOI.

This is where the literature stops and our project begins. A second check after AOI is needed and can recognise false calls, but trained once it fails on later production days; drift and a threshold set once explain this, and methods for both exist and have worked elsewhere. For the second check of AOI results, no study has tested them. We therefore reimplemented the method of [7] with our own code as the baseline. Reading their code, we found that the stored threshold belongs to the last model tried during tuning, not to the best model that is finally used; in our reimplementation the threshold is taken from the best model. On this baseline, the threshold method and retraining become the variables of a simulation, and every setting is compared with [7].

2.2 Summary of the reviewed works

Table 4 compares the studies on AOI false calls, and Table 5 compares the studies on drift and retraining. Table 6 maps the methods of level 8 to the variables of our simulation. Only [7] evaluates on later production periods, and no study on AOI false calls retrains its model (Table 4). Retraining helped in every study that tested it, but none of these studies is on AOI (Table 5).

Table 4: Studies on AOI false calls compared with our project.

	Data	Classifiers	Evaluation	Training
Thielen et al. [5]	not public	kNN, neural network, random forest	time order not stated	no retraining reported
Albayari et al. [6]	Siemens AG	XGBoost, random forest, DNN, kNN	several train–test splits	no retraining reported
Pfab & Rothering [7]	Siemens, public [4]	kNN, RF, balanced RF, XGBoost, AutoML	test set + 5 later periods	once
Our project	Siemens, public [4]	same as [7] (RF, balanced RF)	same as [7]	**once vs. retrained (6 schedules), 3 threshold methods**

Table 5: Studies on drift and retraining compared with our project.

	Domain	Task	Remedy studied	Main finding
Gama et al. [16]	data streams (benchmarks)	classification	retrain when the error rate rises (DDM)	drift can be detected from the error rate alone
Gama et al. [12]	survey	any	windows, forgetting, retraining	a periodically retrained sliding window is the usual baseline
Dal Pozzolo et al. [17]	credit card fraud	binary, <1 % positives	daily update on the newest labelled data	recent data rank frauds better
Bauer et al. [13]	weld quality (images)	binary	none (effect of drift)	all metrics drop under drift
Gao et al. [18]	semiconductor materials	regression	fixed schedule, window type	frequency matters most
Our project	SMT, AOI (Siemens)	binary, ≈ 1 % positives	threshold method and retraining schedule	tested in P3

3 Discussion

3.1 Why the earlier results look better than they are

Most studies evaluate on a random or stratified split of data from the same period [5,6]. Under drift, such a split mixes past and future and hides the loss in performance [12]. The evaluation in time order of [7] is therefore the one that matches the use in a factory, and it is the reason why we build on this study. The measure matters as well: with about 1 % true defects, accuracy says little, and recall and specificity must be reported for the class that matters [11].

Table 6: From the literature to the variables of our simulation (the two next steps named in [7]).

Variable	Option	What it does	Sources
Threshold	as in [7], corrected	mean of the five cross-validation thresholds, taken from the best model	[7]
	safety margin	1 % bound held with probability $1-\delta$; needs enough left-out true defects	[14]
	re-set per period	follows a changed share of true defects, also without labels	[15]
Retraining	never	train once, as in [7]	[7]
	on a fixed schedule	per period, weekly, every 3 days, daily; on all data so far or on a recent window	[12, 17, 18]
	on drift	retrain when the error rises	[16]

3.2 Limitations of the reviewed sources

Two of the central sources are preprints [7, 18]; for [7], the published code and run logs allow us to check the results ourselves. The industrial targets of Siemens [3] are targets, not measured results. The manufacturing evidence for retraining [18] concerns regression, not classification, and the drift experiment of [13] works on images. The labels of the Siemens data come from operators and may contain errors [4]. Finally, all evidence for the AOI case comes from one production line; whether the findings transfer to other lines is not known.

4 Conclusion

The literature leads from the general goal of zero-defect manufacturing [1] to one open question. AOI raises many false calls that people check by hand [2, 4], and a second check with machine learning can release many of them [5, 6]. On later production days, however, it misses 5–15 % of the true defects instead of at most 1 % [7]. Drift [12, 13] and a threshold set once [14, 15] explain this failure, and retraining has helped in related fields [17, 18], but for AOI no study has tested either remedy.

Position of our project. In their discussion, Pfab and Rothering name as "potential next steps to stabilize the modeling" the "improving [of] the method for setting the threshold and handling the distribution drifts within the dataset" [7], without testing either. Our project focuses on exactly these two steps. We do not propose a new method; we test whether the standard remedies work for this real problem. Our reimplementation of [7], with the threshold taken from the best model, is the baseline. On it, a design of experiments varies the threshold method (as in [7] corrected, a safety margin [14], re-set per period [15]), the retraining schedule (never, per period, weekly, every 3 days, daily, on drift [16]), the retraining data (all data so far or a recent window [12]) and the model (random forest or balanced random forest), each setting with the ten seeds of [7]. The question is: how do missed true defects and released false calls on later production periods change when threshold and retraining change, compared with training once? Every result is shown next to the published result of [7]; its targets (at most 1 % missed true defects, at least 40 % fewer manual checks) serve as reference lines. In the end, the simulation becomes a small software tool in which threshold and retraining are settings. Whatever the result, it answers a question that the original study left open.

References

[1] F. Psarommatis, G. May, P.-A. Dreyfus, and D. Kiritsis, "Zero defect manufacturing: state-of-the-art review, shortcomings and future directions in research," *Int. J. Production Research*, vol. 58, no. 1, pp. 1–17, 2020, doi: 10.1080/00207543.2019.1605228.

[2] A. A. R. M. Abu Ebayyeh and A. Mousavi, "A review and analysis of automatic optical inspection and quality monitoring methods in electronics industry," *IEEE Access*, vol. 8, pp. 183192–183271, 2020, doi: 10.1109/ACCESS.2020.3029127.

[3] Siemens, "Siemens and NVIDIA advance AI-powered PCB inspection from Digital Twin to factory," press release, Plano, TX, USA, Oct. 21, 2026. [Online]. Available: `https://news.siemens.com/en-us/siemens-nvidia-pcba-digital-twin-factory/`

[4] K. Pfab, R. Eichler, A. Mallandur, and M. Rothering, "Data of automated optical inspection of surface-mounted technology electronic production," *Data in Brief*, vol. 53, Art. no. 110110, Apr. 2024, doi: 10.1016/j.dib.2024.110110. Data: Mendeley Data, v2, doi: 10.17632/99jzmh9658.2.

[5] N. Thielen, D. Werner, K. Schmidt, R. Seidel, A. Reinhardt, and J. Franke, "A machine learning based approach to detect false calls in SMT manufacturing," in *Proc. 43rd Int. Spring Seminar on Electronics Technology (ISSE)*, Demanovska Valley, Slovakia, 2020, pp. 1–6, doi: 10.1109/ISSE49702.2020.9121044.

[6] J. Albayari, P. Patra, D. Patel, and D. Santos, "Enhancing false call detection in automated optical inspection using XGBoost classifier and XAI algorithms," in *Flexible Automation and Intelligent Manufacturing (FAIM 2025)*, Lecture Notes in Mechanical Engineering. Cham, Switzerland: Springer, 2026, pp. 509–530, doi: 10.1007/978-3-032-05610-8_51.

[7] K. Pfab and M. Rothering, "Towards improved research methodologies for industrial AI: a case study of false call reduction," arXiv:2506.14521 [cs.LG], Jun. 2025. Source code and run logs: `https://gitlab.com/research9841542/FCR_1`. Open-access copy: PDF in our repository.

[8] L. Grinsztajn, E. Oyallon, and G. Varoquaux, "Why do tree-based models still outperform deep learning on typical tabular data?" in *Advances in Neural Information Processing Systems 35 (NeurIPS 2022), Datasets and Benchmarks Track*, 2022. Open-access copy: PDF in our repository.

[9] L. Breiman, "Random forests," *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001, doi: 10.1023/A:1010933404324.

[10] C. Chen, A. Liaw, and L. Breiman, "Using random forest to learn imbalanced data," Dept. of Statistics, University of California, Berkeley, Tech. Rep. 666, 2004. Open-access copy: PDF in our repository.

[11] H. He and E. A. Garcia, "Learning from imbalanced data," *IEEE Trans. Knowledge and Data Engineering*, vol. 21, no. 9, pp. 1263–1284, Sep. 2009, doi: 10.1109/TKDE.2008.239.

[12] J. Gama, I. Žliobaitė, A. Bifet, M. Pechenizkiy, and A. Bouchachia, "A survey on concept drift adaptation," *ACM Computing Surveys*, vol. 46, no. 4, Art. no. 44, pp. 1–37, 2014, doi: 10.1145/2523813.

[13] J. C. Bauer, S. Trattnig, F. Vieltorf, and R. Daub, "Handling data drift in deep learning-based quality monitoring: evaluating calibration methods using the example of friction stir welding," *J. Intelligent Manufacturing*, vol. 37, no. 2, pp. 759–774, 2025, doi: 10.1007/s10845-025-02569-6.

[14] X. Tong, Y. Feng, and J. J. Li, "Neyman-Pearson classification algorithms and NP receiver operating characteristics," *Science Advances*, vol. 4, no. 2, Art. no. eaao1659, 2018, doi: 10.1126/sciadv.aao1659. Open-access copy: PDF in our repository.

[15] M. Saerens, P. Latinne, and C. Decaestecker, "Adjusting the outputs of a classifier to new a priori probabilities: a simple procedure," *Neural Computation*, vol. 14, no. 1, pp. 21–41, 2002, doi: 10.1162/089976602753284446.

[16] J. Gama, P. Medas, G. Castillo, and P. Rodrigues, "Learning with drift detection," in *Advances in Artificial Intelligence – SBIA 2004*, Lecture Notes in Computer Science, vol. 3171. Berlin, Germany: Springer, 2004, pp. 286–295, doi: 10.1007/978-3-540-28645-5_29. Open-access copy: PDF in our repository.

[17] A. Dal Pozzolo, G. Boracchi, O. Caelen, C. Alippi, and G. Bontempi, "Credit card fraud detection: a realistic modeling and a novel learning strategy," *IEEE Trans. Neural Networks and Learning Systems*, vol. 29, no. 8, pp. 3784–3797, 2018, doi: 10.1109/TNNLS.2017.2736643. Open-access copy: PDF in our repository.

[18] M. Gao, J. M. Perathoner, A. L. Bonin, S. Eulig, and G. Klesse, "Robust and reliable AI for predictive quality in semiconductor materials manufacturing with MLOps and uncertainty quantification," arXiv:2605.07752 [cs.LG], 2026. Open-access copy: PDF in our repository.